

基于PXE的物理机 Kubernetes集群管理方案

关于作者

- 工程师, 熟悉Python以及Go语言
- Docker早期用户, 容器架构实践者, 开源爱好者
- 曾在多次技术活动中分享容器相关经验
- 开源软件资讯站 Porter.io 作者
- 目前在经营 KubeUp.com , 致力于协助客户公司以Kubernetes为基础结合众多开源组件进行底层架构, 打造基于容器技术的高效开发运维流程, 同时为企业使用Kubernetes过程遇到的问题提供技术支持和服务



twitter: [@adieu](https://twitter.com/adieu)
github: github.com/adieu
website: www.adieu.me

您平常如何创建Kubernetes集群？

手动创建

选择OS

- CentOS
- Ubuntu
- CoreOS



创建VM

- VMWare
- OpenStack
- Cloud



初始化

- kubeadm
- bootkube
- Manual

自动创建



手动创建物理机集群

选择OS

- CentOS
- Ubuntu
- CoreOS



手动安装

- Manual



初始化

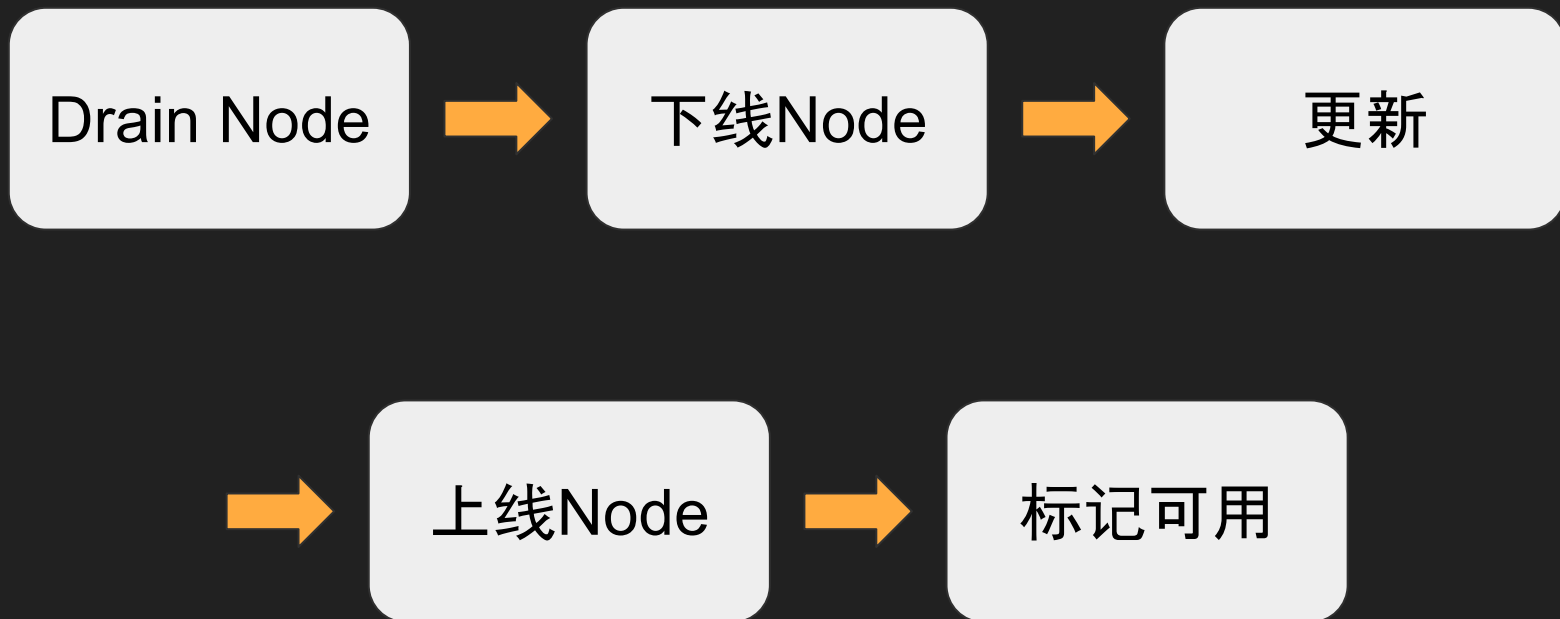
- kubeadm
- bootkube
- Manual

自动创建物理机集群

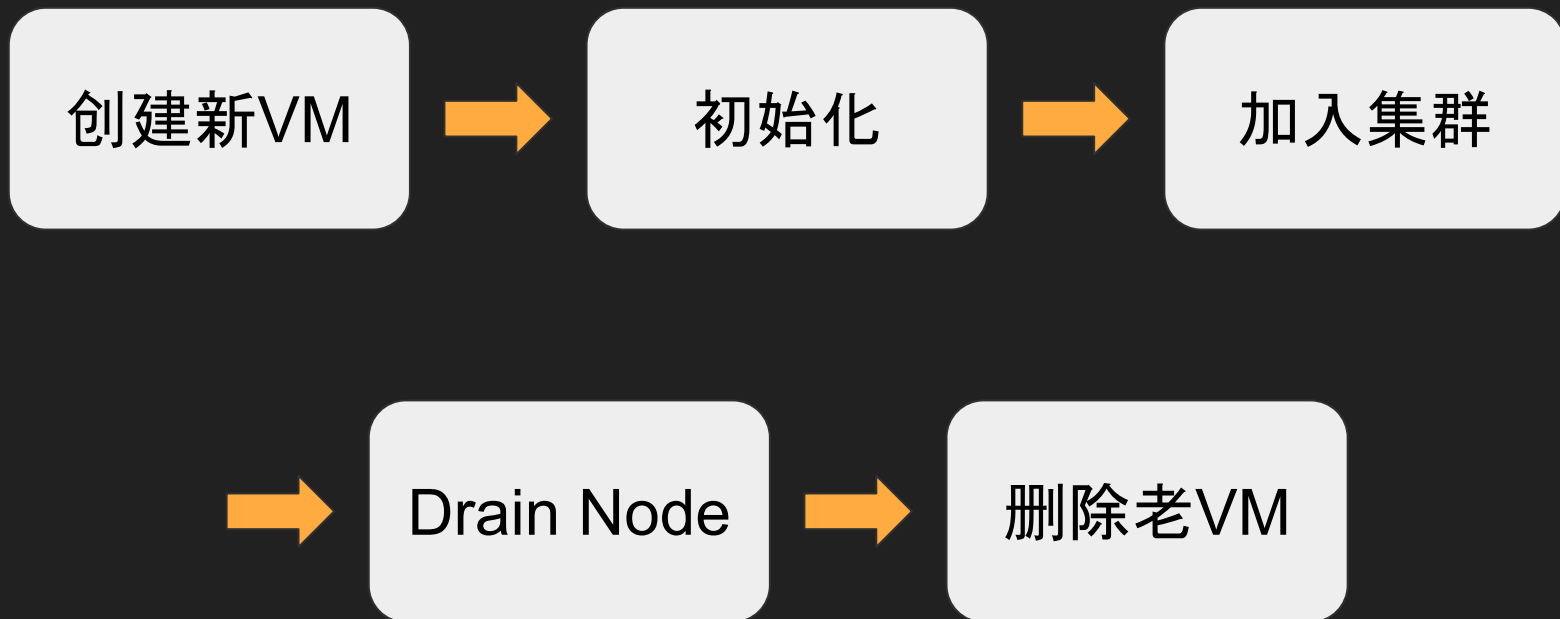


您平常如何升级Kubernetes集群？


手动升级



手动升级



手动升级物理机集群



自动升级物理机集群



如何自动化运维物理机集群？

OpenStack



Q SEARCH

SOFTWARE ▾

USERS ▾

COMMUNITY ▾

MARKETPLACE

EVENTS ▾

LEARN ▾

DOCS

JOIN ▾

LOG IN

From Boston to your browser, stream OpenStack Summit sessions

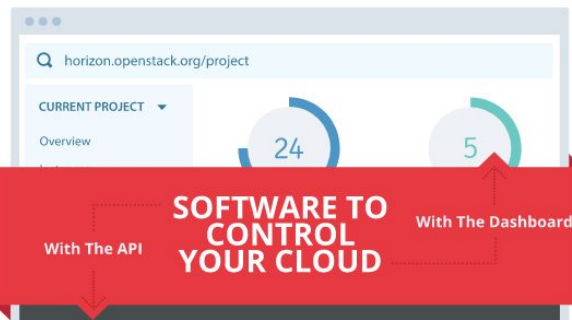
Browse through use cases, demos and announcements from the OpenStack Summit Boston this week.

Watch now >


Open source software for creating private and public clouds.

OpenStack software controls large pools of compute, storage, and networking resources throughout a datacenter, managed through a [dashboard](#) or via the [OpenStack API](#). OpenStack works with [popular enterprise and open source technologies](#) making it ideal for heterogeneous infrastructure.


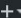

[Hundreds of the world's largest brands](#) rely on OpenStack to run their businesses every day, reducing costs and helping them move faster. OpenStack has a strong ecosystem, and users seeking commercial



将OpenStack运行在Kubernetes之上

 This repository Search

[Pull requests](#) [Issues](#) [Marketplace](#) [Gist](#)

openstack / openstack-helm

[Watch](#) 20 [Unstar](#) 33 [Fork](#) 17

[Code](#) [Pull requests](#) 0 [Projects](#) 0 [Insights](#)

Helm charts for deploying OpenStack on Kubernetes <http://openstack.org>

[798 commits](#) [1 branch](#) [1 release](#) [34](#)




Branch: master [New pull request](#) [Create new file](#)

Jenkins committed with openstack-gerrit Merge "Fix ceph instructions in multinode guide."

barbican	Refactor start scripts, to support graceful termination
bootstrap	Remove executable flag from Chart.yaml
ceph	Remove executable flag from Chart.yaml
cinder	Merge "Refactor start scripts, to support graceful termination"
doc/source	Fix ceph instructions in multinode guide.
etcd	Update deployment for kubernetes 1.6
glance	Helm Test: Glance
heat	Refactor start scripts, to support graceful termination
helm-toolkit	Helm Test: Framework & Keystone
horizon	Refactor start scripts, to support graceful termination
ingress	Update deployment for kubernetes 1.6

This repository Search

[Pull requests](#) [Issues](#) [Marketplace](#) [Gist](#)

stackanetes / stackanetes

[Watch](#) 30 [Star](#) 185 [Fork](#) 42

[Code](#) [Issues](#) 4 [Pull requests](#) 7 [Projects](#) 0 [Insights](#)

Deploy and manage OpenStack on Kubernetes

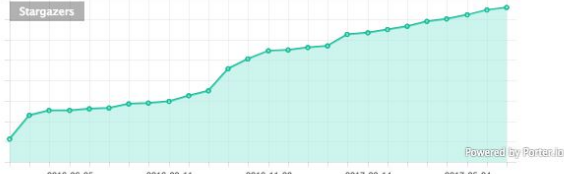
[556 commits](#) [6 branches](#) [3 releases](#) [14 contributors](#) [Apache-2.0](#)

Branch: master [New pull request](#) [Create new file](#) [Upload files](#) [Find file](#) [Clone or download](#)

PiotrProkop committed on GitHub Merge pull request #153 from Quentin-M/mem_usage Latest commit 6ecb457 on Mar 14

cinder	Rename deployment.replicas by deployment.control_replicas	6 months ago
elasticsearch	Remove privileged=true and uid=0 on most services.	8 months ago
etcd	Changing label from rabbit to stackanetes	8 months ago
glance	Enforce 'replicas: 1' on glance-api when local storage is used	5 months ago

Stargazers



Powered by Portento

News 1

OpenStack on Kubernetes
Oct 26 2016

PXE拯救世界



WIKIPEDIA
The Free Encyclopedia

[Main page](#)
[Contents](#)
[Featured content](#)
[Current events](#)
[Random article](#)
[Donate to Wikipedia](#)
[Wikipedia store](#)

Interaction

[Help](#)
[About Wikipedia](#)
[Community portal](#)
[Recent changes](#)
[Contact page](#)

Tools

[What links here](#)
[Related changes](#)
[Upload file](#)
[Special pages](#)
[Permanent link](#)
[Page information](#)
[Wikidata item](#)
[Cite this page](#)

[Print/export](#)

[Create a book](#)

Not logged in [Talk](#) [Contributions](#) [Create account](#) [Log in](#)

Article [Talk](#)

[Read](#)

[Edit](#)

[View history](#)



Preboot Execution Environment

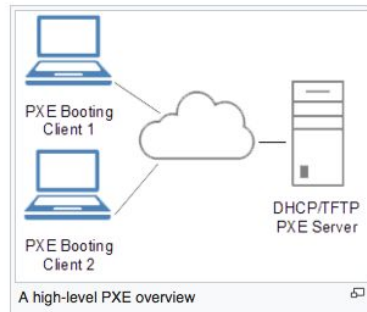
From Wikipedia, the free encyclopedia

In [computing](#), the **Preboot eXecution Environment** (**PXE**, sometimes pronounced as *pixie*^[1]) specification describes a standardized [client-server](#) environment that [boots](#) a software assembly, retrieved from a network, on PXE-enabled clients. On the client side it requires only a PXE-capable [network interface controller](#) (NIC), and uses a small set of industry-standard network protocols such as [DHCP](#) and [TFTP](#).

The concept behind the PXE originated in the early days of protocols like [BOOTP](#)/[DHCP](#)/[TFTP](#), and as of 2015 it forms part of the [Unified Extensible Firmware Interface](#) (UEFI) standard. In modern data centers, PXE is the most frequent choice^[2] for [operating system](#) booting, installation and deployment.

Contents [hide]

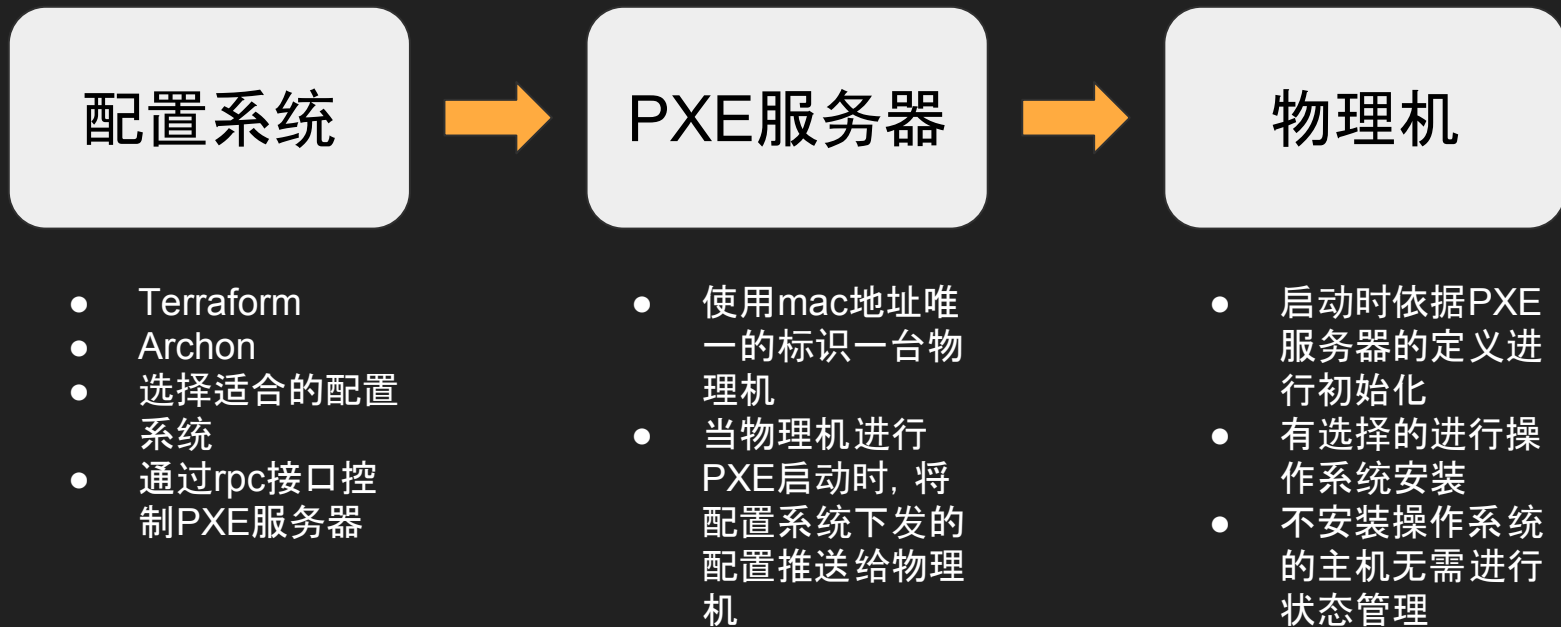
- [Overview](#)
- [Details](#)
- [Integration](#)
- [Availability](#)
- [Acceptance](#)
- [Sibling environments](#)
- [Descendant environments](#)
- [IETF standards documentation](#)
- [See also](#)
- [References](#)
- [External links](#)




Overview [edit]

Since the beginning of computer networks, there has been a persistent need for [client](#) systems which can boot appropriate software images, with appropriate configuration parameters, both




架构



Matchbox

 This repository Search

Pull requestsIssuesMarketplaceGist

coreos / matchbox

Watch 59Star 358Fork 111


[Code](#) [Issues 39](#) [Pull requests 5](#) [Insights](#)

Network boot and provision Container Linux clusters (e.g. etcd3, self-hosted Kubernetes, and more)

811 commits6 branches9 releases32 contributorsApache-2.0

Branch: masterNew pull request

Create new fileUpload filesFind fileClone or download

 dghubble committed on GitHub Merge pull request #564 from coreos/remove-cmdline Latest commit 50a3d11 13 hours ago

Documentation	Doc/deployment.md: Cp local config to correct location	5 days ago
cmd	cmd/matchbox: fix -address parsing when built with Go 1.8	3 months ago
contrib	*: Update docs references to v0.6.1	13 days ago
examples	examples/terraform: Configure whether to install CL from cache	7 days ago
matchbox	matchbox: Remove Profile cmdline map field	2 days ago
scripts	CHANGES.md: Prepare for a v0.6.1 docs point release	13 days ago
tests/smoke	examples/terraform: Configure whether to install CL from cache	7 days ago
vendor	glide.yaml: Update and vendor the crypto openpgp package	13 days ago
.dockerignore	Update dockerignore to slim the build context	5 months ago
.gitignore	contrib/rpm: Fixing the source URL format	28 days ago

CoreOS

[Products](#) ▾[Open Source](#) ▾[Resources](#)[Company](#) ▾[Contact Us](#)[Sign Up](#)

A container-focused OS that's designed for painless management in large clusters

[Overview](#)[Documentation](#)[Release Notes](#)

CoreOS Container Linux 1437.0.0 Documentation

[1437.0.0 \(latest\)](#) ▾

These docs should cover the entire lifecycle of your Container Linux machines. Anyone can submit changes to [these docs via GitHub](#). For more in-depth support, jumping into [#coreos on IRC](#), [emailing the dev list](#) or [filing a bug](#) are recommended.



Running Container Linux


Container Linux runs on most cloud providers, virtualization platforms and bare metal servers. Running a local VM on your laptop is a great dev environment. Following the [Quick Start guide](#) is the fastest way to get set up.

Ignition




```
storage:
  {{ if index . "pxe" }}
  disks:
    - device: /dev/sda
      wipe_table: true
      partitions:
        - label: ROOT
  filesystems:
    - name: root
      mount:
        device: "/dev/sda1"
        format: "ext4"
        create:
          force: true
          options:
            - "-LR00T"
  {{end}}
```


```
systemd:
  units:
    - name: docker.service
      enable: true
    - name: locksmithd.service
      dropins:
        - name: 40-etcd-lock.conf
          contents: |
            [Service]
            Environment="REBOOT_STRATEGY=etcd-lock"
            Environment="LOCKSMITHD_ETCD_CAFILE=/etc/ssl/etcd/etcd-ca.crt"
            Environment="LOCKSMITHD_ETCD_CERTFILE=/etc/ssl/etcd/etcd-client.crt"
            Environment="LOCKSMITHD_ETCD_KEYFILE=/etc/ssl/etcd/etcd-client.key"
            Environment="LOCKSMITHD_ENDPOINT={{.etcd_endpoints}}"
    - name: kubelet.path
      enable: true
      contents: |
        [Unit]
        Description=Watch for kubeconfig
        [Path]
        PathExists=/etc/kubernetes/kubeconfig
        [Install]
        WantedBy=multi-user.target
```

Archon

 This repository Search

Pull requestsIssuesMarketplaceGist



 kubeup / archon

Unwatch 6Unstar 76Fork 8

[Code](#) [Issues 6](#) [Pull requests 0](#) [Projects 0](#) [Wiki](#) [Settings](#) [Insights](#)


Cluster operation the Kubernetes way [Edit](#)

[kubernetes](#) [go](#) [container](#) [Manage topics](#)

111 commits2 branches6 releases2 contributorsApache-2.0

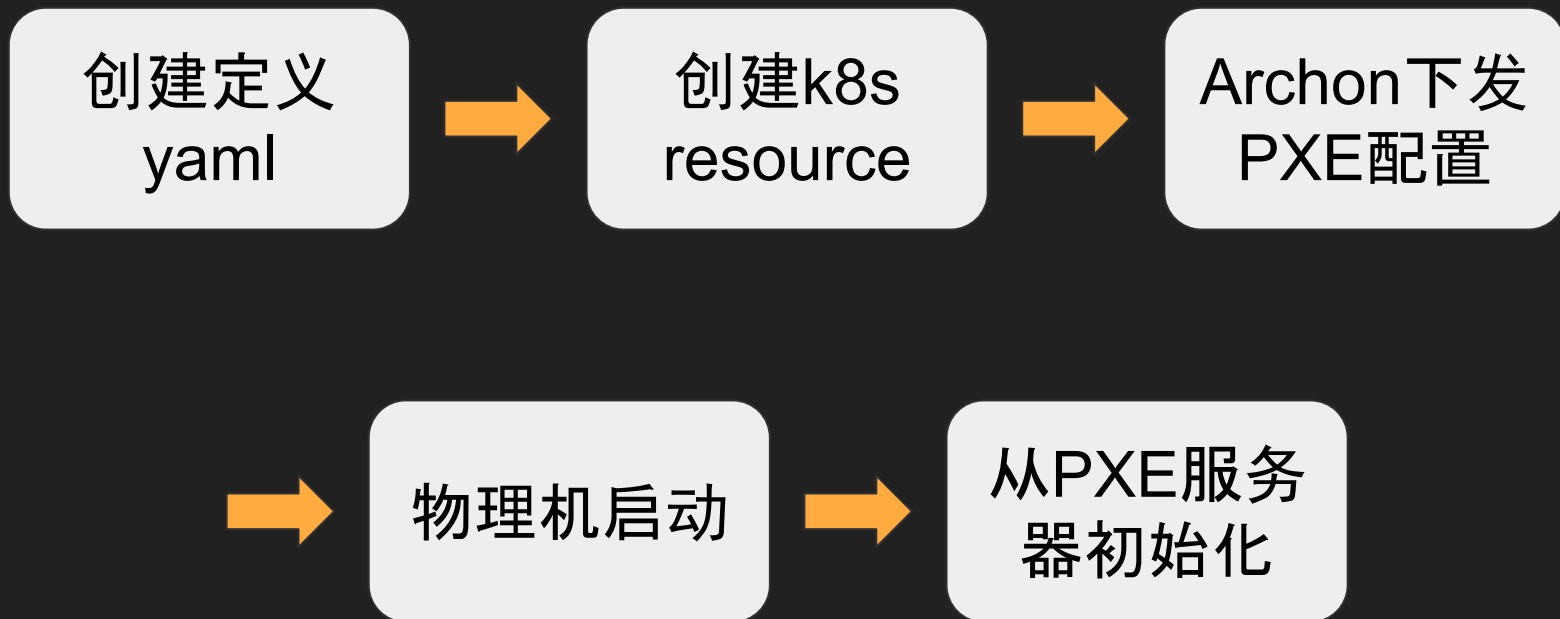
Branch: master New pull request

Create new fileUpload filesFind fileClone or download

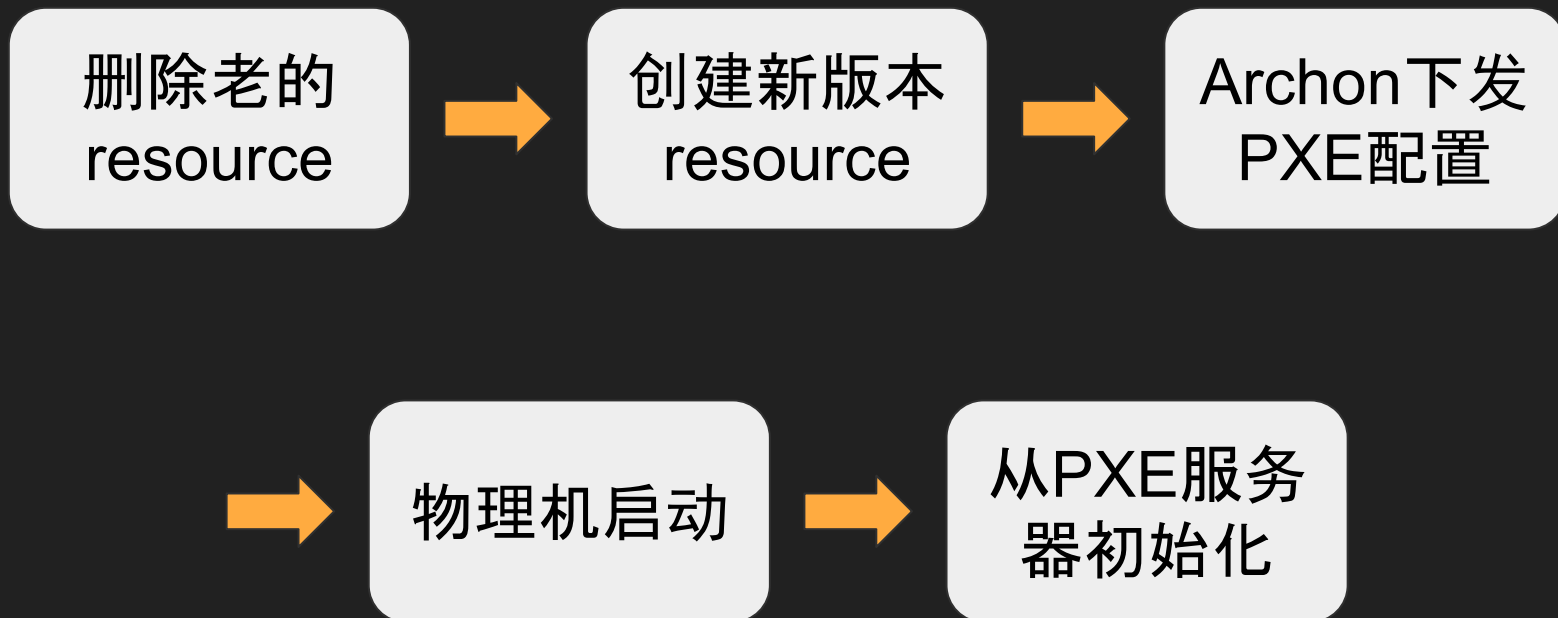
 ledzep2 committed with adieu fix #7 Latest commit 2838f51 20 days ago

cmd/archon-controller	use k8s 1.6	3 months ago
docs	Add aliyun example.	4 months ago
example	aliyun centos example and code	25 days ago
pkg	fix #7	15 days ago
vendor	aliyun centos example and code	25 days ago
.gitignore	initial commit	6 months ago
Dockerfile	update docker base image. add --test-run option	3 months ago
LICENSE	minor fixes of docs	5 months ago
Makefile	update docker base image. add --test-run option	3 months ago


自动创建物理机集群






自动升级物理机集群



terraform-provider-matchbox

 This repository Search Pull requests Issues Marketplace Gist

 17  23  4

[Code](#) [Issues 1](#) [Pull requests 0](#) [Insights](#)

Terraform provider for on-premise / bare-metal via matchbox

38 commits

1 branch


3 releases

2 contributors

Apache-2.0


Branch: master New pull request

Create new file Upload files Find file Clone or download




 mcuadros committed with dghubble matchbox: Add Profile raw_ignition support Latest commit 011ced7 16 days ago

examples	examples: Move examples into upstream matchbox examples	2 months ago
matchbox	matchbox: Add Profile raw_ignition support	14 days ago
scripts	Add Makefile release target for linux and darwin	2 months ago
vendor	glide.yaml: Vendor matchbox/tlsutil subpackage	14 days ago
.gitignore	Add Makefile release target for linux and darwin	2 months ago
.travis.yml	Add .travis.yml file for Travis CI	2 months ago
CONTRIBUTING.md	LICENSE: Add initial license and notices	2 months ago
DCO	LICENSE: Add initial license and notices	2 months ago
LICENSE	LICENSE: Add initial license and notices	2 months ago
Makefile	Makefile: Fix release target to use cross-compilation flags	24 days ago

Sextant

 This repository Search

Pull requestsIssuesMarketplaceGist

k8sp / sextant

Watch 28Star 61Fork 25


[Code](#) [Issues 33](#) [Pull requests 1](#) [Projects 1](#) [Wiki](#) [Insights](#)

Fully automatic installation of CoreOS+Kubernetes clusters

985 commits8 branches3 releases14 contributorsApache-2.0

Branch: develop New pull request

Create new fileUpload filesFind fileClone or download

 typhoonzero committed on GitHub Merge pull request #536 from k8sp/fix_image_name

Latest commit 239ff7a on May 3

doc	Add Install CentOS design	7 months ago
docker	fix 519	3 months ago
golang	fix download kernel logits.	2 months ago
install-ceph	Merge branch 'test_cc_centos' of https://github.com/k8sp/sextant into...	6 months ago
logo	add logo	7 months ago
scripts	fix image save name and README	a month ago
testdata	Add testdata/example.yaml and test check_prerequisites	8 months ago
vm-cluster	Remove docker version	5 months ago
.gitignore	put build ops in docker	8 months ago
.travis.yml	Add .travis.yml	11 months ago
LICENSE.txt	Rephrase copyright holders. Taking https://github.com/tensorflow/tens...	9 months ago

Demo

进一步改进计划

- 基于IPMI(Intelligent Platform Management Interface)实现远程控制物理机启动顺序以及重启物理机
- 与DCIM(Data Center Infrastructure Management)系统集成, 从它同步物理机基础数据, 将物理机加入被管理列表
- 全套方案完全开源, 欢迎感兴趣的朋友一起来完善

总结

- 在物理机集群环境中, 缺少有效的远程控制手段, 传统运维方法中手动比例较高
- PXE可以用来自动化安装和配置物理机操作系统
- 辅以配套的管理工具, 可以实现从配置定义到配置下发一整套的物理机自动化运维方案
- 将不可变基础架构的管理思路拓展到物理机, 可以将运维流程标准化和自动化, 但是新的运维流程将非常依赖配置管理工具
- 使用场景很关键, 需要有一定的集群规模才能平衡引入新系统带来的管理负担的增加

Q & A